

Distributed Systems, Summer Term 2019

Exercise Sheet 3

In all the exercises of this problem set, we consider the synchronous message passing model on a graph, where nodes operate in synchronous rounds and all nodes start a computation together at time 0. We assume that initially, nodes do not know the IDs of their neighbors.

1. Leader Election in General Graphs

Consider the following leader election algorithm. For simplicity, we assume that every node knows the graph diameter D . Every node u stores the largest ID it has seen in variable x_u . Each node $u \in V$ carries out the following algorithm.

```
Node  $u$  initially sets  $x_u := \text{ID}(u)$  and sends  $x_u$  to its neighbors.  
for  $D - 1$  rounds do  
  if  $x_v > x_u$  for the largest value  $x_v$  that  $u$  received then  
     $u$  sets  $x_u := x_v$  and sends  $x_u$  to all neighbors from which it has not received a value  
    equal to  $x_v$ .
```

After D rounds, the value x_u of each node u equals the largest ID in the network.

What is the message complexity of this algorithm (in terms of n)? Give an example that shows that your given bound is asymptotically tight.

Sample Solution

In every round of the algorithm, at most $2m$ messages are sent (at most one in each direction via each edge). The message complexity is therefore $O(D \cdot m) = O(n^3)$. For a given n we construct a graph in the following way: Nodes with IDs $n, \dots, \lceil n/2 \rceil$ form a path, nodes with IDs $\lceil n/2 \rceil, \dots, 1$ form a clique. Let v be the node with ID $\lceil n/2 \rceil$. For $\lceil n/2 \rceil - 1$ rounds, v updates x_v and forwards it to its $\lceil n/2 \rceil - 1$ neighbors in the clique, which in the following round update their values and forward them to $\lceil n/2 \rceil - 2$ nodes (every node in the clique except themselves and v). So for $\lceil n/2 \rceil - 1$ rounds, there are at least $(\lceil n/2 \rceil - 1)(\lceil n/2 \rceil - 2)$ messages sent per round.

2. Leader Election via Radius Growth

We generalize the radius growth algorithm for leader election from the lecture to arbitrary graphs. Assume that every node knows the number of nodes in the graph. The algorithm consists of phases $i = 0, 1, 2, \dots$. Let C_i be the set of leader candidates at the beginning of phase i . Set $C_0 = V$ (initially, each node is a leader candidate). Phase i of the algorithm consists of 2^i rounds. The algorithm terminates when $2^i \geq n$. In phase i , each node $u \in V$ carries out the following algorithm.

If $u \in C_i$, u initializes $x_u := \text{ID}(u)$ and sends x_u to its neighbors (otherwise, u initializes $x_u := -1$).
for $2^i - 1$ rounds **do**
 if $x_v > x_u$ for the largest value x_v that u received **then**
 u sets $x_u := x_v$ and sends x_u to all neighbors from which it has not received a value
 equal to x_v .
if $u \in C_i \wedge x_u == \text{ID}(u)$ **then**
 u joins C_{i+1} (i.e., u stays a candidate)

- Show that the number of messages sent in phase i is $O(\min\{2^i, |C_i|\} \cdot m)$.
- Show that $|C_i| \leq \frac{4n}{2^i}$ for each phase i .
- Show with a) and b) that the message complexity of the algorithm is at most $O(m\sqrt{n} \log n)$.
- For $m = \Omega(n^2)$, the upper bound from c) becomes $O(n^{5/2} \log n)$. Give an example network on which the algorithm requires $\Omega(n^{5/2})$ messages.

Sample Solution

- Phase i consists of 2^i rounds and in one round at most $2m$ messages are sent, so there are at most $2^i \cdot m$ messages. On the other hand, in each phase only the values of nodes in $|C_i|$ are broadcasted and each edge transports such a value at most once in each direction which yields an upper bound of $O(|C_i| \cdot m)$ messages.
- In phase i , two candidates $u, v \in C_i$ have distance at least 2^{i-1} , as otherwise the one with the smaller ID would have lost its candidate status at the end of phase $i - 1$. It follows that $B_{2^{i-2}}(u)$ and $B_{2^{i-2}}(v)$ (the balls with radius 2^{i-2} around u and v) are disjoint. If $2^{i-2} \leq D/2$ (i.e., $i \leq \log(D)+2$), then $B_{2^{i-2}}(u)$ contains at least 2^{i-2} nodes. It follows that $n \geq |C_i| \cdot 2^{i-2}$.
- There are at most $\log n$ phases. For phases $i = 0, \dots, \log(n)/2$ we have $2^i \leq \sqrt{n}$. For phases $i = \log(n)/2, \dots, \log n$ we have $|C_i| \leq 4\sqrt{n}$. With a) it follows that in each phase at most $O(m \cdot \sqrt{n})$ messages are sent.
- For a given n , we construct a graph with n nodes in the following way. For $i = 1, \dots, \sqrt{n}$, there is a line of length i : $(v_1^1), (v_1^2, v_2^2), (v_1^3, v_2^3, v_3^3), \dots, (v_1^{\sqrt{n}}, v_2^{\sqrt{n}}, \dots, v_{\sqrt{n}}^{\sqrt{n}})$. Nodes v_i^i form a clique together with $n - \sum_{i=1}^{\sqrt{n}} i = \Omega(n)$ additional nodes. For each i , node $v_1^{\sqrt{n}-i}$ has the i^{th} -largest ID in the graph.

When running the algorithm on this graph, there are at least \sqrt{n} rounds in which an ID reaches the clique which is larger than the ones it has seen before. Then all $\Omega(n)$ nodes v in the clique update their value x_v and send it to their $\Omega(n)$ neighbors in the following round. So there are at least \sqrt{n} rounds with a message complexity of $\Omega(n^2)$.

3. Leader Election in Complete Graphs

In a complete graph, one can trivially solve leader election in one round if every node sends its ID to all its neighbors. This requires $\Omega(n^2)$ messages. The following algorithm uses less messages at the cost of a slightly higher time complexity.

The algorithm consists of phases $i = 1, 2, \dots$. Let C_i be the set of leader candidates at the beginning of phase i . Set $C_1 = V$ (initially, each node is a leader candidate). In phase i , each node $u \in V$ carries out the following algorithm.

if $u \in C_i$ **then**

u sends a probe message containing its ID to $\min\{2^i, n-1\}$ arbitrary neighbors.

Let v be the node with the largest ID from which u received a probe message

if $\text{ID}(v) > \text{ID}(u)$ **then**

u sends back an acknowledgement to v

if u received 2^i acknowledgements **then**

u joins C_{i+1} (i.e., u remains a candidate)

- a) Argue that the algorithm solves leader election and analyze its time complexity.
- b) Show that $|C_i| \leq \frac{n}{2^{i-1}}$ for each phase $i \geq 1$.
- c) Analyze the message complexity.

Sample Solution

- a) The node with maximum ID remains candidate in every phase. In phase $\lceil \log n \rceil$ all candidates send probes to all other nodes and thus the node with maximum ID is the only one surviving. As a phase consists of two rounds, the runtime is $O(\log n)$.
- b) For $u \in C_i$, let $A_u \subseteq V$ be the set of nodes from which u received an acknowledgement. We have $|A_u| = 2^{i-1}$ and for $u, v \in C_i$ we have $A_u \cap A_v = \emptyset$. Therefore, $n \geq |C_i|2^{i-1}$.
- c) In phase i , there are at most $2 \cdot |C_i| \cdot 2^i$ messages sent. With b) it follows that the message complexity of each phase is $O(n)$. As the algorithm has $\log n$ phases, its message complexity is $O(n \log n)$.