

Distributed Systems, Summer Term 2019

Exercise Sheet 6

1 Leader Election with Failures

Consider the leader election problem on a complete graph in the synchronous message passing model. That is, every node has a unique ID and at end of the algorithm, every node that did not crash has to output the ID of the leader node.

Explain how to adapt the $(f+1)$ -round lower bound proof for consensus from the lecture to show that if at most $f \leq n - 2$ processes may fail during the protocol, at least $f + 1$ rounds are needed to solve leader election.

Sample Solution

In the $(f+1)$ -round lower bound proof for consensus we built a similarity chain between an execution with f rounds with all inputs 0 and no crashes and an execution with f rounds with all inputs 1 and no crashes. But in fact the concrete input values did not play any role when building this chain, which means that we proved an even stronger property: We can build a similarity chain between two executions with f rounds and no crash failures and arbitrary input values.

Now assume there is a protocol for solving leader election in $f \leq n - 2$ rounds. Take pairwise distinct integers $x_1, \dots, x_n, y_1, \dots, y_n$ which represent the IDs of two leader election instances. Let E be the execution with inputs x_1, \dots, x_n and no crash failures and E' the execution with inputs y_1, \dots, y_n and no crash failures. As there is a similarity chain between E and E' , all nodes must have the same output in both executions which is a value among x_1, \dots, x_n . But then in E' the output is not one of the IDs.

Alternative proof: Assume there is an algorithm solving leader election in at most f rounds, i.e., after f rounds all non crashed nodes output the same value which is the ID of some node. Now assume every node v has an input $x_v \in \{0, 1\}$. Node v changes its ID to $ID'(v) = 2 \cdot ID(v) + x_v$ (take the binary representation of $ID(v)$ and concatenate it with x_v). Then run the leader election algorithm, but every node only outputs the last bit of the ID of the leader node. Then we have solved binary consensus which we know is unsolvable in f rounds.

2 k -set agreement

A generalization of consensus is the k -set agreement problem: Every node has an input value and at the end every node has to output a value such that the following properties are fulfilled:

1. Agreement: There must not be more than k different output values.
2. Validity: Every node must output a value which was input of some node.

Show that on a complete graph in the synchronous message passing model with at most f node failures, the k -set agreement problem is solvable in $\lfloor f/k \rfloor + 1$ rounds.

Sample Solution

Algorithm: Each node v maintains a value x_v which is initially set to v 's input. For $\lfloor f/k \rfloor + 1$ rounds, each node executes the following code

```
Broadcast  $x_v$   
Receive values  $y_1, \dots, y_m$   
Update  $x_v$  to  $\min\{x_v, y_1, \dots, y_m\}$ 
```

Analysis: Validity is clear. For agreement assume there is a round with less than k crash failures. Let x_1, \dots, x_ℓ with $\ell < k$ be the values of the nodes that crashed in the current round and let V be the set of nodes that survive the round. Then at the end of the round, each node in V chooses a value among $x_1, \dots, x_\ell, \min\{x_w | w \in V\}$ as the values in $\{x_w | w \in V\}$ arrive at every node. So every node chooses among $\ell + 1 \leq k$ different values and the number of different values can only be decreased in the following rounds. This means that as soon as there is a round with less than k crashes, the algorithm fulfills agreement. In $\lfloor f/k \rfloor + 1$ rounds, there must be at least one round with less than k crashes.