

Information Retrieval

WS 2017 / 2018

Lecture 5, Tuesday November 21st, 2017
(Fuzzy Search, Edit Distance, q-Gram Index)

Prof. Dr. Hannah Bast
Chair of Algorithms and Data Structures
Department of Computer Science
University of Freiburg

Overview of this lecture

■ Organizational

- Experiences with ES4 Compression, Codes, Entropy
- Date of written exam **Mon Feb 19 or Tue Feb 20 ?**

■ Contents

- Fuzzy search type breifurg, find freiburg
- Edit Distance a standard similarity measure
- Q-gram Index index for efficient fuzzy search

ES5: implement error-tolerant prefix search using a q-gram index and prefix edit distance

We have compiled a beautiful new dataset for you:
2.9M entities from Wikidata with scores and descriptions

■ Summary / excerpts

- None of you found the exercises easy at first glance
- But some of you sat down and eventually figured it out and realized that in retrospect it wasn't that hard

"Looking back now the exercises were actually easy. But I think that's always the case when doing math exercises :)"
- Some of you gave up or did only a part of the exercises

"I am not a fan of this exercise sheet"
- First time use of LaTeX for some ... **don't worry, it's worth it!**
- In Exercise 3, confusing to use L_i for the length of the i -th inverted list, when we used it for the i -th code length before

- Favorite movies watched in 2017 (selection)
 - One night in Paris ... "why? because of the romance"
 - Toy Story 2 ... "because it's very funny"
 - Captain Fantastic ... "not the type of movie I usually watch, but it made me reconsider"
 - Blade Runner 2049 ... "I totally expected Hollywood to ruin it but I was pleasantly surprised"
 - La La Land ... several mentions
 - The Hobbit ... "the scene with Gollum (not Golomb) was great"
 - Fack ju Göhte 3 ... "I really enjoy feeling my own IQ decrease constantly while watching movies like this one. Maybe that's the reason why this sheet took me so long to complete."

Experiences with ES4 3/3

$$\lceil x \rceil \leq x + 1$$

■ Proof sketch for Exercise 4.2

$$L_i = \lfloor \frac{i}{M} \rfloor + 1 + \lceil \log_2 M \rceil$$

$$M = \lceil \frac{\ln 2}{p} \rceil \leq \frac{\ln 2}{p} + 1$$

- Show that "Gollum" encoding with modulus $M = \lceil \ln 2/p \rceil$ is optimal if the prob. for number i is $p_i = (1-p)^{i-1} \cdot p$

we need to show: $L_i \leq \log_2 \frac{1}{p_i} + O(1)$

$$(I) \log_2 \frac{1}{p_i} = \log_2 \frac{1}{(1-p)^{i-1} \cdot p} = (i-1) \cdot \log_2 \frac{1}{1-p} + \log_2 \frac{1}{p}$$

$$(II) L_i \leq \frac{i \cdot p}{\ln 2} + 1 + \log_2 \left(\frac{\ln 2}{p} + 1 \right) + 1$$

$$= \frac{(i-1)p}{\ln 2} + \frac{p}{\ln 2}$$

$$\leq 2$$

$$\leq \frac{(i-1)p}{\ln 2} + \log_2 \frac{1}{p} + 4 + \underbrace{\log_2 (3 \ln 2)}_{\leq 2}$$

$$\leq (i-1) \cdot \log_2 \frac{1}{1-p} + \log_2 \frac{1}{p} + 6$$

$$\leq \log_2 \frac{1}{p_i} + 6$$

$$\ln 2 = 0.69 \dots$$

$$2 \cdot \ln 2 \geq 1$$

$$\frac{2 \cdot \ln 2}{p} \geq \frac{1}{p} \geq 1$$

Hint:

$$1+x \leq e^x \quad \checkmark x$$

$$1-p \leq e^{-p}$$

$$\frac{1}{1-p} \geq e^p$$

$$\ln \frac{1}{1-p} \geq p$$

$$\log_2 \frac{1}{1-p} = \frac{\ln \frac{1}{1-p}}{\ln 2} \geq \frac{p}{\ln 2}$$

■ Problem setting

- Given a "dictionary" = a list of "names" of any kind

For ES5: a list of **2,920,180** entities from Wikidata

- For a given query, find matching names from that dict.

Query: frei Match: freiburg

prefix search

Query: fr*rg Match: freiburg

wildcard search

Query: breifurg Match: freiburg

fuzzy search

- Similar challenges as for our search so far:

Challenge 1: good model of what **matches**

Challenge 2: preprocess the input (= build a suitable index), so that we find the matching names **fast**

■ Possible origins for the dictionary

- Popular queries extracted from a query log

Basis for Google's query-suggestion feature

- Words + common phrases from a text collection

Extracting common phrases from a given text collection is an interesting problem by itself, however, not one we will deal with in this course

- A list of names of entities

For example: person names, movie titles, places, street addresses, ...

■ Combining matching and search

- One could simply search for the top match, for example:

Type: freib Search: freiburg

- Or one could search for several matches

Type: freib Search: freiburg OR freibach OR ... OR ...

- In today's lecture, we will only look at the problem of finding matching names in a list of names

The search part is also interesting when the number of matching strings is very large; then a simple OR of a lot of strings will be too slow and we need better solutions

■ Simple solution

- Iterate over all strings in the dictionary, and for each check whether it matches
- This is what the Linux commands **grep** and **agrep** do

```
grep -x uni.* <file>
```

```
grep -x un.*ity <file>
```

```
agrep -x -2 univerty <file>
```

All matching lines in **<file>** will be output

The option **-x** means match whole line (not just a part)

The option **-2** means allow up to two "errors" ... next slide

- Simple solution, check match of single string

- Given a query q and a string s

- **Prefix search:** **easy-peasy**

- Just compare q and the first $|q|$ characters of s ... can be accelerated by finding the first match with a binary search

- **Wildcard search:** **also easy if only one ***

- If $q = q_1 * q_2$, check that $|s| > |q_1| + |q_2|$ and then compare the first $|q_1|$ characters of s with q_1 and the last $|q_2|$ characters of s with q_2

- **Fuzzy search:** **more complicated**

- Compute edit distance between q and s ... slides 12 – 17

- Simple solution, time complexity
 - The time complexity is obviously $n \cdot T$, where
 - n = #records, T = time for checking a single string
 - For fuzzy search, $T \approx 1\mu s$... find out yourself in ES5
 - In search, we always want interactive query times
 - Respond times feel interactive until about **100ms**
 - So the simple solution is fine for up to $\approx 100K$ records
 - For larger input sets, we need to pre-compute something
 - We will build a **q-gram index** ... slides 18 – 25

Edit distance 1/6

Vladimir
Levenshtein
1935 - 2017



UNI
FREIBURG

■ Definition ... aka Levenshtein distance, from 1965

- Definition: for two strings x and y

$ED(x, y) :=$ minimal number of tra'fo's to get from x to y

- Transformations allowed are:

$insert(i, c)$: insert character c at position i

$delete(i)$: delete character at position i

$replace(i, c)$: replace character at position i by c

$x = \text{DOOF}$
 $\text{DOOF} \rightarrow \text{BOOF} \quad \text{REPLACE}(1, B)$
 $\text{BOOF} \rightarrow \text{BLOOF} \quad \text{INSERT}(2, L)$
 $\text{BLOOF} \rightarrow \text{BLOEF} \quad \text{REPLACE}(4, E)$
 $\text{BLOEF} \rightarrow \text{BLOED} \quad \text{REPLACE}(5, D)$
 $y = \text{BLOED}$

IMPORTANT :

This only proves

$$ED(x, y) \leq 4$$

Edit distance 2/6

■ Some simple notation

- The empty word is denoted by ε
- The length (#characters) of x is denoted by $|x|$
- Substrings of x are denoted by $x[i..j]$, where $1 \leq i \leq j \leq |x|$

■ Some simple properties

- $ED(x, y) = ED(y, x)$
- $ED(x, \varepsilon) = |x|$
- $ED(x, y) \geq \text{abs}(|x| - |y|)$ $\text{abs}(z) = z \geq 0 ? z : -z$
- $ED(x, y) \leq ED(x[1..n-1], y[1..m-1]) + 1$ $n = |x|, m = |y|$

DOOF BLOED DOO BLOE
=4 =3

Edit distance 3/6

NOTE:

Recursive
implementation
of this is **NOT**
a good idea
(takes $3^{l \times l}$ time)

■ Recursive formula

- For $|x| > 0$ and $|y| > 0$, $ED(x, y)$ is the minimum of

$$(1a) \quad ED(x[1..n], y[1..m-1]) + 1$$

$$(1b) \quad ED(x[1..n-1], y[1..m]) + 1$$

$$(1c) \quad ED(x[1..n-1], y[1..m-1]) + 1 \quad \text{if } x[n] \neq y[m]$$

$$(2) \quad ED(x[1..n-1], y[1..m-1]) \quad \text{if } x[n] = y[m]$$

- For $|x| = 0$ we have $ED(x, y) = |y|$

- For $|y| = 0$ we have $ED(x, y) = |x|$

For a proof of that formula, see e.g. "Algorithmen und Datenstrukturen" SS 2017, Lecture 11a, slides 14 – 20

Edit distance 4/6

■ Algorithm for computing $ED(x, y)$

- The recursive formula from the previous slide naturally leads to the following dynamic programming algorithm
- Takes time and space $\Theta(|x| \cdot |y|)$

Handwritten DP table for $ED(x, y)$ where $x = \text{DOOF}$ and $y = \text{BLOED}$.

			$\overset{y}{\rightarrow}$	ε	B	L	O	E	D	
ε		0	1	2	3	4	5			$ED(\varepsilon, \text{BLOED})$
\times	D	1	1	2	3	4	4			
\downarrow	O	2	2	2	2	3	4			
	O	3	3	3	2	3	4			
	F	4	4	4	3	3	4			$ED(\text{DOOF}, \text{BLOED})$

■ Prefix edit distance

- The prefix edit distance between x and y is defined as

$PED(x, y) = \min_{y'} ED(x, y')$ where y' is a prefix of y

- For example

$PED(\text{uni}, \underline{\text{university}}) = 0$... but $ED = 7$

$PED(\text{uniwer}, \underline{\text{university}}) = 1$... but $ED = 5$

- Important for fuzzy search-as-you type suggestions

By now, all the large web search engines have this feature, because it is so convenient for usability

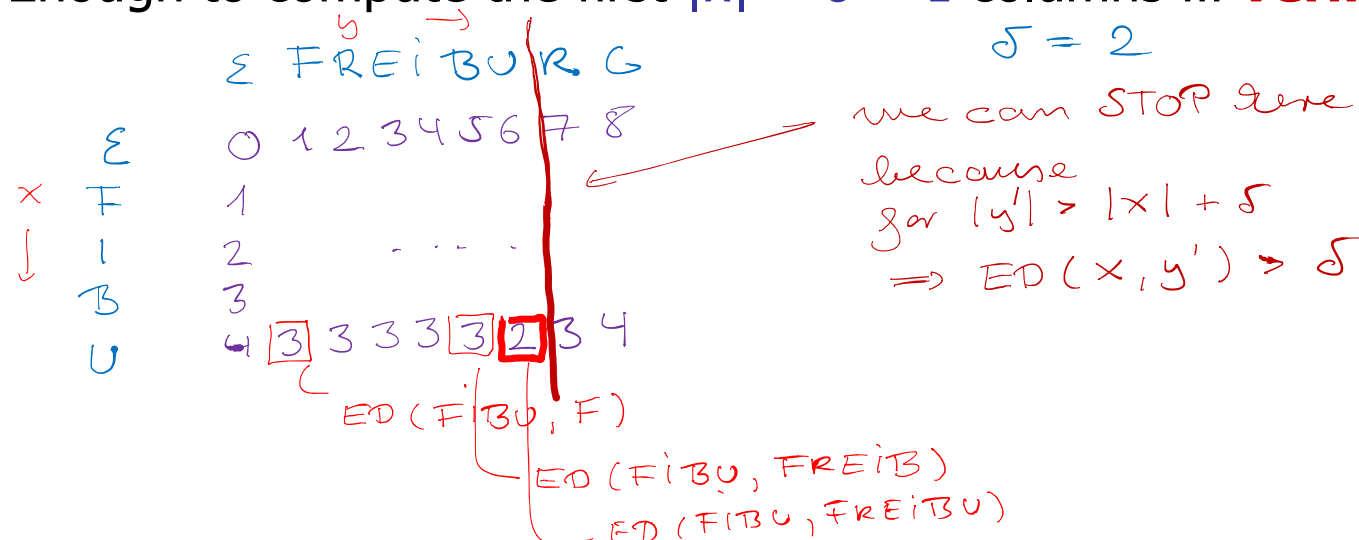
Edit distance 6/6

$$\text{PED}(\text{FIBU}, \text{FREIBURG}) = 2$$

■ Computation of the PED

- Compute the entries of the $|x| \cdot |y|$ table, just as for ED
- The PED is just the minimum of the entries in the last row
- Important optimization: when $|x| \ll |y|$ and you only want to know if $\text{PED}(x, y) \leq \delta$ for some given δ :

Enough to compute the first $|x| + \delta + 1$ columns ... **verify !**



q-Gram Index 1/8

■ Definition of a q-gram

- The q-grams of a string are simply all substrings of length q

freiburg: fre, rei, eib, ibu, bur, urg

$q = 3$

- The number of q-grams of a string x is exactly $|x| - q + 1$

$8 - 3 + 1 = 6$

- For fuzzy search, we will **pad** the string with $q - 1$ special symbols (we use \$) in the beginning and in the end

freiburg \rightarrow \$\$freiburg\$\$

3-grams: \$\$f, \$fr, fre, rei, eib, ibu, bur, urg, rg\$, g\$\$

- The number of q-grams is then $|x| + q - 1$, where x is the original string (the padding adds $q - 1$ q-grams per side)

We will see in a minute, why that padding is useful

■ Definition of a q-gram index

- For each **q**-gram store an inverted list of the strings (from the input set) containing it, sorted lexicographically

\$fr : **fr**aberg, **fr**allach, **fr**eiburg, **fr**eiberg, **fr**ouville, ...

ibu : **bi**burg, **fr**eiburg, gar**ci**bu**eu**y, sei**bu**ttendorf, ...

As usual, store **ids** of the strings, not the strings themselves

- A q-gram index for a collection of words is also an inverted index, just with different objects than a document index:

document index : lists of **doc ids**, one for each **word**

q-gram index : lists of **words ids**, one for each **q-gram**

- Let's now adapt our code from Lecture 1 to q-grams ...

■ Space consumption

- Each record x contributes $|x| + Q$ ids to the inverted lists, where Q is the total number of padding characters

For ES5, we have one-sided padding and $q=3$, hence $Q=2$

- The **total number of ids** in the inverted lists is hence:

$$\#chars + Q \cdot \#words = \#chars \cdot (1 + Q / AVWL)$$

- With 4 bytes per id, $Q = 2$ and $AVWL = 8$, the **total size** of the inverted lists in bytes would hence be $5 \cdot \#chars$

Note: for 1 byte per char, $\#chars$ is the size of the input

- This can be reduced significantly using **compression**

For ES5, it is fine to store the lists uncompressed

q-Gram Index 4/8

■ Fuzzy search with a q-gram index, using ED ... part 1

- Consider x and y with $ED(x, y) = \delta$
- Intuitively: if x and y are not too short, and δ is not too large, they will have one or more q -grams in common
- Example: $x = \text{HILLARY}$, $y = \text{HILARI}$, $q = 3$, $\delta = 2$
 $\$\$ \text{HILLARY} \$\$ \rightarrow \$\$ \underline{\text{H}}, \$\$ \underline{\text{HI}}, \underline{\text{HIL}}, \text{ILL}, \text{LLA}, \underline{\text{LAR}}, \text{ARY}, \text{RY}\$, \text{Y}\$ \$$
 $\$\$ \text{HILARI} \$\$ \rightarrow \$\$ \underline{\text{H}}, \$\$ \underline{\text{HI}}, \underline{\text{HIL}}, \text{ILA}, \underline{\text{LAR}}, \text{ARI}, \text{RI}\$, \text{I}\$ \$$

Number of q -grams in common: 4

- Note: the padding in the beginning gives us two additional 3-grams in common here (because no mistake in first letter)

The more q -grams in common, the more efficient the query algorithm on slide 22 will work

■ Fuzzy search with a q-gram index, using ED ... part 2

- Formally: let x' and y' be the padded versions of x and y

Then: $\text{comm}(x', y') \geq \max(|x|, |y|) - 1 - (\delta - 1) \cdot q$

Example from slide before: $|x| = 7$, $|y| = 6$, $\delta = 2$, $q = 3$

Hence: $\text{comm}(x', y') \geq \max(7, 6) - 1 - 1 \cdot 3 = 3$

In the example, actually: $\text{comm}(x', y') = 4$

Verify: in the worst case, $\text{comm}(x', y') = 3$ can happen

- **Proof:** consider the longer string, which has $\max(|x|, |y|) + q - 1$ q-grams ... because of the left and right \$ padding

Then one tra'fo (insert / delete / replace) change at most q q-grams, and hence δ tra'fos change at most $\delta \cdot q$ q-grams

■ Fuzzy search with a q-gram index, using ED ... part 3

- Given a query x and a q -gram index for the input strings
- Compute q -grams of x' and fetch their inverted lists

For example: $x = \text{HILARI}$, $x' = \$\$ \text{HILARI} \$\$$

Fetch lists for: $\$ \$ H$, $\$ HI$, HIL , ILA , LAR , ARI , $RI \$$, $I \$ \$$

- Merge these lists and for each word in the union keep a **count** of the number of lists in which it is contained, for example:

HILLARY:	4	(contains $\$ \$ H$, $\$ HI$, HIL , LAR)
HAEMOPHILIA:	2	(contains $\$ \$ H$ and HIL)
SOLAR:	1	(contains only LAR)

This step considers each word that contains at least one of the q -grams, which is usually many more than actually match

■ Fuzzy search with a q-gram index, using ED ... part 4

- For each record y in the merge results, check whether the count is $\geq \max(|x|, |y|) - 1 - (\delta - 1) \cdot q$

NO: discard this y , we know that $ED(x, y) > \delta$

YES: compute $ED(x, y)$ and check if $ED(x, y) \leq \delta$

- Let's continue our example

$x = \text{HILARY}$, $q = 3$, $\delta = 2$... hence $-1 - (\delta - 1) \cdot q = -4$

$y_1 = \text{HILLARY}$: 4 $\max(|x|, |y_1|) - 4 = 3 \rightarrow \text{YES}$

$y_2 = \text{HAEMOPHILIA}$: 2 $\max(|x|, |y_2|) - 4 = 6 \rightarrow \text{NO}$

$y_3 = \text{SOLAR}$: 1 $\max(|x|, |y_3|) - 4 = 2 \rightarrow \text{NO}$

- So for this example, we only have to compute $ED(\text{HILARY}, \text{HILLARY})$... which is 2, hence HILLARY is output as a match

- Changes when using the **PED** ... needed for ES5
 - We use the same algorithm, but with a different bound
 - Assume that $\text{PED}(x, y) \leq \delta$
 - Let x' and y' be x and y with $q - 1$ times \$ to the **left only** (padding on the right makes no sense for prefix search)
 - Then we have: $\text{comm}(x', y') \geq |x| - q \cdot \delta$

Note that for $\delta = 1$, this is ≥ 1 only for $|x| > q$

- **Proof:** Consider x , which has exactly $|x| - q + 1$ q-grams

Then one tra'fo (insert / delete / replace) changes at most q q-grams, and hence δ tra'fos change at most $\delta \cdot q$ q-gram

References

- Textbook

 - Section 3: Tolerant Retrieval, in particular:

 - Section 3.2: Wildcard queries

 - Section 3.3: Spelling correction

- Wikipedia

 - <http://en.wikipedia.org/wiki/N-gram>

 - [http://en.wikipedia.org/wiki/Approximate string matching](http://en.wikipedia.org/wiki/Approximate_string_matching)

 - [http://en.wikipedia.org/wiki/Levenshtein distance](http://en.wikipedia.org/wiki/Levenshtein_distance)