

# Information Retrieval

WS 2017 / 2018

Lecture 9, Tuesday December 19<sup>th</sup>, 2017  
(Clustering, K-Means)

Prof. Dr. Hannah Bast  
Chair of Algorithms and Data Structures  
Department of Computer Science  
University of Freiburg

# Overview of this lecture

---

## ■ Organizational

- Your experiences with ES8      Vector Space Model
- Christmas break, 2 weeks      [U+1F600](#) [U+1F36A](#)

## ■ Contents

- Clustering      Definition and example
- K-Means      Algorithm and analysis
- K-Means for text      All kinds of practical advice

ES9: a few pencil-and-paper exercises which should  
test and deepen your understanding

(this is the kind of tasks you might also get in the exam)

## ■ Summary / excerpts

- Conceptually easy, but the devil is in the details

"I like this sheet: not too hard, but one needs to be careful"

- Tricky to figure out how to do things efficiently in scipy

"realized quickly that todense method is evil"

- At least one of you used Java and regretted it (linear algebra code is so much more convenient in Python)

- Many of you still catching up on your programming practice

"I feel like I am learning more and more about Python and Java programming ... thank you, I realized that I really need more training ... I am having fun (I guess)."

## ■ Results

- Three score variants : BM25 and tf
- Three normalizations : rows, columns, none

- Understand that:

Row normalization is similar (not identical) to idf

Column normalization is similar (not identical) to the document length normalization of BM25

- The best results were achieved with BM25

Which is more evidence (based on very small benchmark though) that the BM25 normalizations are indeed good

## ■ Informal definition

- Given elements  $x_1, \dots, x_n$  from a **metric space**

Metric space = there is a measure of distance between any two elements from the space

- Group the elements into clusters  $C_1, \dots, C_k$  such that

**Intra**-cluster distances are as small as possible

**Inter**-cluster distances are as large as possible


We will make this notion more precise on slide 7

- We assume that  $k$  (the number of clusters) is given as part of the input

# Clustering 2/3


$k = 3$

## ■ Example



# Clustering 3/3

one  
cluster



## ■ Centroids and RSS

- Assume we have a **centroid**  $\mu_i$  for each cluster  $C_i$

Intuitively, a centroid is a single element from the metric space that "represents" the cluster

- Let  $c_i$  be the index of the cluster to which  $x_i$  is assigned

Each element belongs to one cluster = **hard** clustering,  
in Lecture 10 we will also see soft clustering

- Then we define the **residual sum of squares** as

$$\text{RSS} = \sum_{i=1, \dots, k} \sum_{x \in C_i} (x - \mu_i)^2 = \sum_{i=1, \dots, n} (x_i - \mu_{c_i})^2$$

$$(x_i - \mu)^2 = |x_i - \mu|^2$$

The sum of the squares of all intra-cluster distances

## ■ Algorithm

- Basic idea: find a local optimum of the **RSS** by greedily minimizing it in every step
- Initialization: pick a set of centroids
  - For example, pick a random subset from the input
- Then alternate between the following two steps
  - (A)** Assign each element to its nearest centroid
  - (B)** compute new centroids as average of elems assigned to it
- Let's first look at [a nice demo](#) and then show that both steps can only **decrease** the RSS



## ■ Step A (assign to nearest centroid)

- Recall:  $RSS = \sum_{i=1, \dots, n} (x_i - \mu_{c_i})^2$
- In Step A, the centroids  $\mu_1, \dots, \mu_k$  are fixed and we want to find those  $c_1, \dots, c_n$  that minimize the RSS:

$$\min_{c_1, \dots, c_n} \sum_{i=1, \dots, n} (x_i - \mu_{c_i})^2 = \sum_{i=1, \dots, n} \min_{c_i} (x_i - \mu_{c_i})^2$$

Each summand can be minimized independently

- $\min_{c_i} (x_i - \mu_{c_i})^2 = \min_{c_i} |x_i - \mu_{c_i}|$

The square distance is min. when the distance is min.

- $|x_i - \mu_{c_i}|$  is minimized for  $c_i = \operatorname{argmin}_j |x_i - \mu_j|$

In words: by assigning  $x_i$  to its nearest centroid

## ■ Step B (recompute centroids)

- Recall:  $RSS = \sum_{i=1, \dots, k} \sum_{x \in C_i} (x - \mu_i)^2$
- In Step B, the clusters  $C_1, \dots, C_k$  are fixed and we want to find the centroids  $\mu_1, \dots, \mu_k$  that minimize the RSS:

$$\min_{\mu_1, \dots, \mu_k} \overbrace{\sum_{i=1, \dots, k} \sum_{x \in C_i} (x - \mu_i)^2}^{RSS} = \sum_{i=1, \dots, k} \min_{\mu_i} \sum_{x \in C_i} (x - \mu_i)^2$$

- Each  $\min_{\mu_i} \sum_{x \in C_i} (x - \mu_i)^2$  can be solved independently, and we can do that using simple calculus:

$$g(\mu) = \sum_{x \in C} (x - \mu)^2$$

$$g'(\mu^*) = -2 \sum_{x \in C} (x - \mu^*) \stackrel{!}{=} 0$$

$$\Leftrightarrow \underbrace{\sum_{x \in C} x}_{= |C| \cdot \mu} - \underbrace{\sum_{x \in C} \mu^*}_{= |C| \cdot \mu^*} = 0 \quad \Leftrightarrow \mu^* = \frac{1}{|C|} \cdot \sum_{x \in C} x$$

$$g''(\mu) = 2 \sum_{x \in C} 1 = 2|C| > 0 \Rightarrow \text{min at } \mu^*$$

concave  
function

# K-Means 4/9

## ■ Convergence to local RSS minimum

- By what we have just proven, RSS stays equal or decreases in every step (A) and every step (B)
- There are only finitely many clusterings ... think: how many?
- Therefore, the algorithm will terminate if we avoid that it cycles forever between different clusterings with equal RSS
- Solution: deterministic tie breaking in the centroid assignment, when two centroids are equally close


For ES9 prefer the centroid with larger index

Some termination conditions don't need a tie-breaking rule, for example, the last condition from slide 13


$$\begin{aligned} & k_2 \text{ clusters} \\ & n \text{ elements} \\ & \# \text{ clusterings} \\ & = \underbrace{k_2 \cdot \dots \cdot k_2}_{n \text{ times}} = k_2^n \end{aligned}$$

# K-Means 5/9

- A local **RSS** minimum is not always a global one



$$\begin{aligned} RSS &= 2 \cdot \left( \left( \frac{4}{3} \right)^2 + \left( \frac{1}{3} \right)^2 + \left( \frac{5}{3} \right)^2 \right) \\ &= 2 \cdot \frac{16 + 1 + 25}{9} \\ &= 2 \cdot \frac{42}{9} = \frac{84}{9} = 9\frac{1}{3} \end{aligned}$$



$$\begin{aligned} RSS &= 4 \cdot 0.5 + 2 \cdot 0.25 \\ &= 2.5 \end{aligned}$$

Quite a bit smaller !

$$x^2 = 0.5^2 + 0.5^2 = 0.5$$

## ■ Termination condition, options

- **Stop** when no more change in clustering

Optimal, but this can take a **very** long time

- **Stop** after a fixed number of iterations

Easy, but how to guess the right number?

- **Stop** when **RSS** falls below a given threshold

Reasonable, but **RSS** may never fall below that threshold

- **Stop** when decrease in **RSS** falls below a given threshold

Reasonable: we stop when we are close to convergence

For the toy example of ES9.2 take the first condition

## ■ Choice of a good $k$

- **Idea 1:** choose the  $k$  with smallest RSS

Bad idea, because RSS is minimized for  $k = n$

- **Idea 2:** choose the  $k$  with smallest  $RSS + \lambda \cdot k$

Makes sense: RSS becomes smaller as  $k$  becomes larger

But now we have  $\lambda$  as a tuning parameter

Experience shows that for a given kind of application, there is often an input-independent good choice for  $\lambda$ , whereas a good  $k$  depends on the input

For ES9, the number of clusters is always given

When  $RSS = 0$

- When is K-Means a good clustering algorithm

- K-Means tends to produce compact clusters of about equal size

Indeed, it can be shown that K-Means is optimal when the sought for clusters are spherical and of equal size

Whether it's good or not, k-means is used a **lot lot lot** in practice, just because of it's simplicity

## ■ Alternatives

### – **K-Medoids**

Maintain that centroids are elements from the input set

### – **Fuzzy k-means**

Elements can belong to several clusters to varying degrees ... this is sometimes called "soft clustering"

L10 will be about a method for soft clustering (LSI)

### – **EM-Algorithm** (EM = Expectation-Maximization)

General-purpose optimization technique that can also be used for soft clustering



## ■ Representation

- We use the vector space model (VSM) from Lecture 8

Each document = one column of our term-doc matrix

- Centroids are also vectors in this space
- To compute the centroid of a set of documents, just take the average of the document vectors
- Important observation: the document vectors are **sparse**, the centroids become **dense** over time

When implementing k-means, document vectors must be stored in sparse representation, just like in L8

- Note: the term-document matrix can be constructed from an inverted index just as shown in the last lecture

# K-Means for Text Documents 2/6

---

## ■ Running time

- Let  $n$  = #documents,  $m$  = #terms,  $k$  = #clusters
- Assume that each dist computation takes time  $\Theta(D)$
- Then each step (A) takes time  $\Theta(k \cdot n \cdot D)$

Compute **dist** between each documents and each cluster

- Each step (B) takes time  $\Theta(n \cdot m)$

Each of the  $n$  documents is added to one centroid vector,  
and one vector addition takes time  $\Theta(m)$

# K-Means for Text Documents 3/6

$$|x - y| = \sqrt{\sum_i (x_i - y_i)^2}$$

PROOF OF LEMMA:

$$\begin{aligned} |x - y|^2 &= (x - y) \cdot (x - y) \\ &= (x - y) \cdot (x - y) \\ &= \underbrace{x \cdot x}_{|x|^2} + \underbrace{y \cdot y}_{|y|^2} - 2 \cdot x \cdot y \end{aligned}$$

## ■ Distance between two documents

- We use Euclidean distance between the normalized docs:  
 $\text{dist}(x, y) := |x' - y'|$ , where  $x' = x / |x|$  and  $y' = y / |y|$

- Straightforward computation between sparse and dense vector takes time  $\Theta(m)$ , where  $m$  = total number of terms

- **Lemma:**  $|x - y|^2 = |x|^2 + |y|^2 - 2 \cdot x \cdot y$ , where  $x \cdot y$  is the dot product of  $x$  and  $y$

$$\begin{pmatrix} \cdot \\ \cdot \\ \cdot \\ \cdot \\ \cdot \end{pmatrix} \begin{pmatrix} \cdot \\ \cdot \\ \cdot \\ \cdot \\ \cdot \end{pmatrix}$$

- Hence: when  $|x| = |y| = 1$ , then  $\frac{1}{2} \cdot |x - y|^2 = 1 - x \cdot y$

Note: when  $|x| = |y| = 1$ , then  $x \cdot y$  is the cosine of the angle between  $x$  and  $y$  ... a common similarity measure in IR

- Computing dot product for a sparse  $x$  and a dense  $y$  takes only time  $\Theta(M)$ , where  $M$  = number of non-zero entries in  $x$

## ■ Using matrix operations

- Both Steps (A) and (B) can be implemented very efficiently using matrix operations

Some hints and examples on the next two slides

- Use the lemma from the previous slides and make sure that the centroids are  $L_2$ -normalized
- Understand that documents do **not** have to be normalized:

For each document, all we want to know is which centroid is the closest; multiplying that document by a constant factor does not change the relative distances to the centroids

Important take-home message for life: understanding the underlying mathematics can save you **a lot** of work

# K-Means for Text Documents 5/6

## ■ Using matrix operations, Step (A)

- For Step (A), we need to compute the dot products between all documents and all centroids
- Let  $A$  be the term-document matrix (one doc per column)
- Let  $C$  be the term-centroid matrix (one centroid per column)
- Then  $C^T \cdot A$  yields a matrix, where the entry at  $i, j$  is exactly the dot product between centroid  $i$  and document  $j$

$$\begin{matrix} n_i \\ \left( \text{---} \right) \\ C^T \\ k \times m \end{matrix} \cdot \begin{matrix} \left( \begin{matrix} | \\ | \\ | \end{matrix} \right) \\ A \\ m \times n \end{matrix} = \begin{matrix} \left( \begin{matrix} \bullet \\ \bullet \\ \bullet \end{matrix} \right) \\ k \times n \\ \text{all distances between} \\ \text{docs and centroids} \end{matrix}$$

# K-Means for Text Documents 6/6

## ■ Using matrix operations, Step (B)

- For Step (B), we need to **add** the vectors of all documents in the same cluster  $C$ , and then divide by  $|C|$

If one normalizes afterwards, one can drop "divide by  $|C|$ "

- Let  $A$  be the term-document matrix (one doc per column)
- Let  $B$  be a 0-1 matrix where the entry at  $i, j$  is 1 iff document  $i$  is in cluster  $j$
- Then  $A \cdot B$  yields a matrix, where the  $j$ -th column is exactly the sum of all documents assigned to cluster  $j$

$n=5$   
 $k=2$

$$\begin{pmatrix} d_1 & d_2 & d_3 & d_4 & d_5 \\ \vdots & \vdots & \vdots & \vdots & \vdots \end{pmatrix} \cdot \begin{pmatrix} 0 & 1 \\ 0 & 1 \\ 1 & 0 \\ 1 & 0 \\ 0 & 1 \end{pmatrix} = \begin{pmatrix} \vdots & \vdots \\ \mu_1 & \mu_2 \end{pmatrix}$$

$A$   
 $m \times 5$

$B$   
 $5 \times 2$

$\mu_1 = d_3 + d_4$   
 $\mu_2 = d_1 + d_2 + d_5$

# References

---

## ■ Further reading

- Textbook Chapter 16: Flat clustering

<http://nlp.stanford.edu/IR-book/pdf/16flat.pdf>

## ■ Wikipedia

- [http://en.wikipedia.org/wiki/Cluster\\_analysis](http://en.wikipedia.org/wiki/Cluster_analysis)
- <http://en.wikipedia.org/wiki/K-means>
- <http://en.wikipedia.org/wiki/K-medoids>
- [http://en.wikipedia.org/wiki/EM\\_Algorithm](http://en.wikipedia.org/wiki/EM_Algorithm)