

Information Retrieval

WS 2017 / **2018**

Lecture 12, Tuesday January 23rd, 2018
(Knowledge Bases, SPARQL, Translation to SQL)

Prof. Dr. Hannah Bast
Chair of Algorithms and Data Structures
Department of Computer Science
University of Freiburg

Overview of this lecture

Official evaluation:

Tell us in the forum if you did not receive any mail so far

■ Organizational

- Your experiences with ES11
- Exam registration

Naive Bayes

Deadline 31.01.2018

■ Content

- Knowledge bases + SPARQL explanation + examples
- Databases + SQL explanation + examples
- SQLite a lightweight database
- SPARQL to SQL algorithm + example
- Performance joins and join order
- **ES12:** Implement the SPARQL → SQL translation and use it to process SPARQL queries with Python+SQLite

■ Summary / excerpts

- Nice topic / exercise / datasets
- For about half of you it was relatively easy
 - Look at the master solution, it's very little code
- For about half, it took more time due to the usual reasons
 - Finding the right numpy operations took some time, too
- The remaining half had no time to do the exercise

■ Results Genres **H**orror, **D**rama, **D**ocumentary, **C**omedy, **W**estern

- Variant 1: Ho 76%, Dr 79%, Do 77%, Co 65%, We 91%
- Variant 2: Ho 87%, Dr 57%, Do 86%, Co 64%, We 98%
- In Variant 2, all classes are equally frequent in the training set, and hence the class probabilities p_c are all equal

Note that p_c is a factor in the formula for predicting class c

- Some classes have more specific words than others:

Western: gang, town, sheriff, man, ranch, men, father, ...

Comedy: man, life, wife, young, family, money, time, ...

- Excluding stopwords as features does not improve the result quality: in prediction, they contribute equally to all classes

■ Results Ratings

- Variant 1: R 76%, PG-13 30%, PG 36%
- Variant 2: R 52%, PG-13 41%, PG 49%
- Again, having all classes equally frequent in the training data helps the classes which are relatively rare in the original data
- Top words are not really specific for any of the classes

R: life, man, young, find, family, finds, father, wife, ...

PG-13: life, family, man, young, world, find, father, old, ...

PG: life, young, father, man, old, world, family, find, ...

The more specific words come further down in the list (but again, the unspecific words do not really hurt)

■ What is a knowledge base

- A knowledge base is a database of statements about entities and their relations
- Critical: **unique** identifiers for each entity and predicate
- A common format / schema is to express all statements as **subject predicate object** triples:

Nicole Kidman	acted in	Eyes Wide Shut
Brad Pitt	acted in	Burn After Reading
Tom Cruise	acted in	Eyes Wide Shut
Sidney Pollack	acted in	Eyes Wide Shut
Joel Cohen	directed	Burn After Reading
Ethan Cohen	directed	Burn After Reading
Nicole Kidman	married to	Tom Cruise

■ Freebase and WikiData

- Freebase is the largest open general-purpose KB to date

Started by Metaweb in 2007, acquired by Google in 2010

Freebase has become read-only in March 2015 and WikiData has taken over to become **the** standard general-purpose KB

Final size: **3 billion** triples on **60 million** entities

- Wikidata is like Wikipedia for knowledge bases: anybody can contribute (with some amount of editorial control)

Current size: **377 million** triples on **43 million** entities

In Wikidata, entities are called "items" and triples are called "statements"

■ Reification

- Restriction to triples is no real restriction: n-ary relationships can also be represented as triples:

m.0jy6xg	film	Finding Nemo
m.0jy6xg	actor	Ellen DeGeneres
m.0jy6xg	character	Dory
m.0jy6xg	type	Voice

m.0jy6xg is an entity name from Freebase

It's a so-called **mediator** entity, the purpose of which is to serve as a link between the entities it connects

The full Wikidata dataset has a similar mechanism

For simplicity, the dataset for ES12 has no mediators

■ Relation to the "Semantic Web"

- The Semantic Web initiative is concerned with making knowledge base data **explicitly** available on the web
- **Variant 1:** semantic mark-up in normal web pages
Typical format: Microdata or JSON-LD (show example)
- **Variant 2:** web pages containing only structured data
Typical format: RDF (a particular kind of XML)
- No rules that enforce consistent entity or relation names
The hope is that people adhere to standards nevertheless,
and that machines can resolve the remaining heterogeneity
Anyway: this is **not** the topic of this lecture / course

■ What is SPARQL

- The standard query language for knowledge bases

SPARQL = **SPARQL Protocol And RDF Query Language**

- Example query in natural language: actors who are married and played together in at least one movie
- The same query expressed in (simplified) SPARQL

```
SELECT ?person1 ?person2 ?film WHERE {  
  ?person1 acted_in ?film .  
  ?person2 acted_in ?film .  
  ?person1 married_to ?person2  
}
```

■ SPARQL syntax

- In the lecture today, we use a simplified syntax

In "real" SPARQL, names of subjects / predicates / objects may contain whitespace and are surrounded by <...>

- The actual SPARQL syntax is slightly more complicated and has many more features


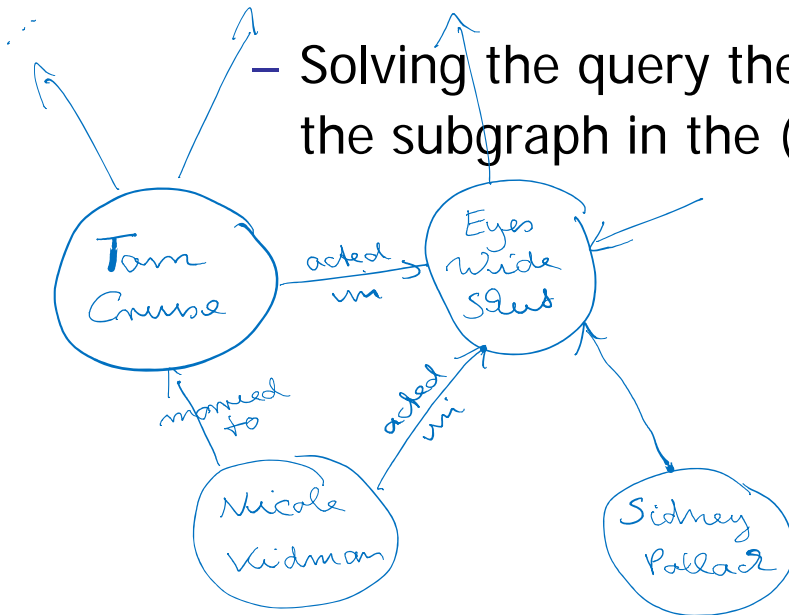
In particular, it involves namespace prefixes, so that names can be made globally unambiguous

See the Wikipedia page or the W3C specification if you are interested

■ SPARQL queries as subgraphs

- One can view a knowledge base as a **graph**, where the nodes are the entities, and the edges are the relations
- A SPARQL query is then a sub-graph with variables at some or all of the nodes
- Solving the query then amounts to finding all matches of the subgraph in the (large) knowledge base graph

KB



■ Introduction

- Data from a knowledge base can also be stored in an ordinary database

This is what we do in the lecture today and for ES12

- The standard query language for databases is SQL

SQL = Structured Query Language

- On the following slides, let us recap the basics from databases and SQL via a few examples

■ What is a database

- For this lecture, a database is a collection of tables, where each table has a fixed number of columns
- For example, we could have one table for each predicate from our knowledge base, with two columns each

Table for "acted in" predicate

actor	film
Brad Pitt	Burn after Reading
Tom Cruise	Eyes Wide Shut
...	...

Table for "married to" predicate

person1	person2
Nicole Kidman	Tom Cruise
Ellen DeGeneres	Portia de Rossi
...	...

For ES12, you should work with **one** big table for the whole database, with three columns (subject, predicate, object)

■ SQL example 1

- Example query FROM **one** table "acted_in" with two columns "actor" and "film"

```
SELECT actor
FROM   acted_in
WHERE  film = "Burn After Reading";
```

In words: all actors from movie "Burn After Reading"

- Principle: select those rows from the specified table which satisfy properties specified in WHERE clause

■ SQL example 2

- Example query FROM **multiple** tables, each with three columns "subject", "predicate", "object"

```
SELECT ex1.subject, ex2.subject, ex1.object
FROM   example AS ex1, example AS ex2
WHERE  ex1.predicate = "acted in"
AND    ex2.predicate = "acted in"
AND    ex1.object = ex2.object
```

In words: all pairs of actors who acted in the same movie

- Principle: selects items from cross-product $T_1 \times \dots \times T_k$ which satisfy properties specified in WHERE clause
- Syntax: use **AS** for unique names of copies of same table; use **table.column** to refer to that column from that table

- A full-fledged database, easy to install and use
 - On Debian/Ubuntu install with: `sudo apt-get install sqlite3`
 - Two types of commands ... examples on next slides
 - SQL commands: must end with a semicolon
 - SQLite commands: start with a dot, no semicolon at end
 - Two modes to start SQLite:
 - `sqlite3` will work on an in-memory database
 - `sqlite3 <name>.db` create database in that file, and if file exists, use database from that file

Let's read our example tables in SQLite using the commands from the next two slides ... it's easy

- Some useful SQLite commands by example
 - Specifies the column separator used for input and output
`.separator " "` use Ctrl+V TAB for TAB !
 - Read table from TSV (tab-separated values) file
`.import film.tsv film`
 - Execute commands from script file (typical suffix is .sql)
`.read <file with commands>`
 - Show execution time of every command
`.timer on`

■ Some useful SQL commands by example

- Create a table with a given schema

```
CREATE TABLE acted_in(actor TEXT, film TEXT);
```

- Create an index for a column of a table

```
CREATE INDEX acted_in_index ON acted_in(actor);
```

- Extract / combine data from tables

```
SELECT * FROM acted_in WHERE ... LIMIT 100;
```

- Delete table / index (without error msg if it's not there)

```
DROP TABLE IF EXISTS acted_in;
```

```
DROP INDEX IF EXISTS acted_in_index;
```

■ Python interface to SQLite

- Executing SQL commands to a SQLite database from within Python is very easy:

```
import sqlite3
db = sqlite3.connect("example.db")
cursor = db.cursor()
cursor.execute("SELECT * FROM table")
for row in cursor.fetchall():
    print("\t".join(row))
```

Beware: the SQLite commands (starting with a dot) cannot be executed from within Python, you need SQLite for those

■ Motivation

- We want to translate a given SPARQL query to a SQL query that gives the desired results on a given database
- In the following example, we use one table per relation

```
CREATE TABLE acted_in(actor TEXT, film TEXT)
```

```
CREATE TABLE married_to(person1 TEXT, person2 TEXT)
```

Note: all elements from one table are from one relation, so we don't need to store the relation name in the table

For ES12, use **one big table** for all the data, with three columns named **subject**, **predicate**, **object**

This is deliberately different from how we did it in the lecture, so that you have to do some thinking yourself

SPARQL to SQL Translation 2/4


■ Example

- SPARQL query

```
SELECT ?p1 ?p2 ?f WHERE {  
  ?p1 acted_in ?f .  
  ?p2 acted_in ?f .  
  ?p1 married_to ?p2      }
```

- SQL query:

```
SELECT  mar.person1, mar.person2, act1.film  
FROM    acted_in AS act1, acted_in AS act2, married_to AS mar  
WHERE   act1.film = act2.film  
AND     act1.actor = mar.person1  
AND     act2.actor = mar.person2
```



■ Algorithm

- It is up to you in ES12, to design a generic algorithm that works for arbitrary basic SPARQL queries

Of the form `SELECT <vars> { <triples> }`

- The algorithm is not difficult, but requires understanding of how the data is stored and how SPARQL and SQL work

It's a perfect exercise to understand the basics !

- On the next slide we give you some valuable advice

■ Algorithm, advice for ES13

- If there are **k** query triples in the **SPARQL** query, have **k** entries in the **FROM** clause of the **SQL** query

FROM freebase as f1, freebase as f2, ... , freebase as fk

- In your code, for each variable from the **SPARQL** query, build an **array** of all its occurrences in the query, e.g.

?x: f1.subject, f2.object, f5.object

- Then, when building the **SQL** query, add the corresponding equalities to the **WHERE** clause, e.g.

WHERE f1.subject = f2.object AND f2.object = f5.object

Note: if **?x** occurs **m** times, **m – 1** equalities are enough

■ Cross product of tables

- Recall that, conceptually, an SQL statement like

`SELECT ... FROM T_1, T_2, \dots, T_k WHERE ...`

selects elements from the **cross-product**


$T_1 \times \dots \times T_k$ (which has $|T_1| \cdot \dots \cdot |T_k|$ elements)

(where some or all of the T_i can be the same table)

■ Joining of tables

- Each ... = ... in the **WHERE clause** effectively ask for a **JOIN** operation between two tables
- Algorithmically, a **JOIN** requires a **list intersection**
- If we **CREATE** an index for the respective tables on the respective join attributes, this list intersection gets fast

E.g., by sorting (a copy of) the table by that attribute



■ Join ordering

- Typical SQL-from-SPARQL queries require multiple joins
- Order of joins can make a **huge** performance difference
- For our example query, the `acted_in` table (actors – films) is about ten times larger than the `married_to` table
- **Join order 1:** look at all pairs of actors who played in the same film, and for each check whether they are married
`materialized all pairs of actors from same film (large)`
- **Join order 2:** look at all married couples and for each get their films and check whether they overlap
`materializes list of films of all married people (small)`

with
all
indexes

■ Join ordering, continued

- Without further ado, **SQLite** seems to take the order of the tables in the **FROM** clause as its join order

```
SELECT married_to.person1, married_to.person2
FROM   acted_in as acted1, acted_in as acted2, married_to
WHERE  married_to.person1 = film1.actor
AND    married_to.person2 = film2.actor
AND    acted1.film = acted2.film;
```

16.4
seconds

Alternatives: (note that there are 6 possible orderings)

```
FROM   married_to, acted_in as acted1, acted_in as acted2
```

2.0
seconds

```
FROM   married_to, acted_in as acted2, acted_in as acted1
```

2.5
seconds

References

■ Textbook

- Nothing in the text book by Manning, Raghavan, Schütze

■ Wikipedia

- http://en.wikipedia.org/wiki/Knowledge_base
- <http://en.wikipedia.org/wiki/SPARQL>
- <http://en.wikipedia.org/wiki/SQL>
- <http://en.wikipedia.org/wiki/SQLite>
- <http://en.wikipedia.org/wiki/Freebase>
- <https://www.wikidata.org>
- https://www.mediawiki.org/wiki/Wikibase/Indexing/RDF_Dump_Format