

# Lecture 16: Clustering

Machine Learning, Summer Term 2019

Michael Tangemann   Frank Hutter   Marius Lindauer

University of Freiburg



# Lecture Overview

1 Motivation

2 Criteria for Clustering


3 K-Means

4 DBSCAN

# What is Cluster Analysis?

Also called: clustering, segmentation analysis, taxonomy analysis, automatic classification, numerical taxonomy, botryology, typological analysis, community detection, ...

(Plot modified from scikit-learn clustering tutorial)




Unsupervised task:

Group objects such that objects within a group are more similar/related to each other (*in some sense*) than to objects of another group.

# What is Cluster Analysis?

Also called: clustering, segmentation analysis, taxonomy analysis, automatic classification, numerical taxonomy, botryology, typological analysis, community detection, ...

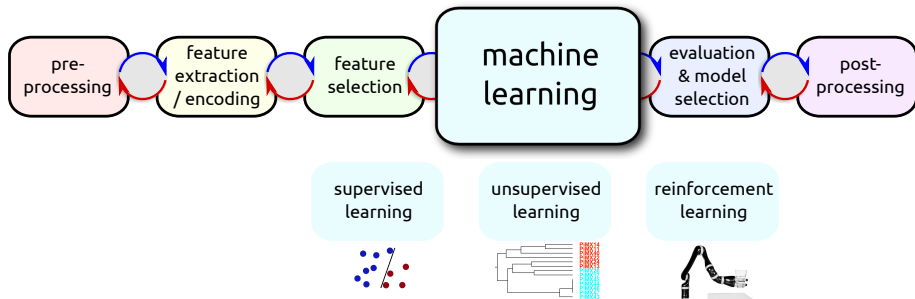
(Plot modified from scikit-learn clustering tutorial)



Unsupervised task:

Group objects such that objects within a group are more similar/related to each other (*in some sense*) than to objects of another group.

# ML Design Cycle



Cluster analysis is an unsupervised learning task

- Ground truth about clusters is not provided → evaluation is tricky!

Given:

- $N$  high dimensional data points  $\mathbf{x}_i \in \mathbb{R}^D$  with  $i = 1 \dots N$ .
- Data is collected in matrix  $\mathbf{X} \in \mathbb{R}^{N \times D}$



# Example Applications (I)

- Medical imaging (fMRI, CT, PET): differentiate between different types of tissues, find tissue boundaries
- Biology: determine communities of organisms in space and time, compute data-driven phylogenetic trees
- Genetics: group DNA sequences into gene families
- Biochemistry / chemistry / pharmacology: group compounds according to their reaction mechanism
- Market research: detect clusters of customers with similar behavior, find market segments
- Social networks: recognize communities
- Search engines: Post-processing of search results into groups of hits that refer to vastly different topics

# Example Applications (II)

- Image segmentation: border detection, track objects
- Anomaly detection: identify outliers in data streams, network attacks, misbehaving software, sensor failures (robotics, production lines), predictive maintenance
- Finance: find stock clusters of similar behaviour
- Text analysis: clustering of documents into topics
- ...




# Lecture Overview

- 1 Motivation
- 2 Criteria for Clustering
- 3 K-Means
- 4 DBSCAN


# How could Clusters be Determined?

Which metrics might be used to define clusters?



# How could Clusters be Determined?

Which metrics might be used to define clusters?



Zoo of clustering methods available, that exploit e.g.:

- distance/similarity function (between cluster members, between members of different clusters)
- connectivity structure using distances → single/avg/max linkage clustering, graph-based → clique
- centroid + neighborhood
- densities
- expected distributions

# Lecture Overview

- 1 Motivation
- 2 Criteria for Clustering
- 3 K-Means**
- 4 DBSCAN

# K-Means Clustering (Steinhaus, 1957)

Find set  $C = C_1, \dots, C_k$  of  $k$  clusters represented by cluster centroids  $\mu_k$  such, that the clusters have equal variance.

→ Minimize the *inertia* or *within-cluster sum-of-squares* criterion:

$$\operatorname{argmin}_C \sum_{i=1}^k \sum_{\mathbf{x} \in C_i} \|\mathbf{x} - \mu_j\|^2$$

Observations:

- Cluster centroids  $\mu_j$  do not need to be points of the training data sets
- **Unfortunately NP-hard problem!**
- Clustering can be represented by Voronoi tessellation

# K-Means Clustering


Practical solution by heuristic approximation

(e.g. Lloyd's algorithm (1957, 1982), similar to expectation-maximization):

Initialize  $k$  data points as cluster centroids.

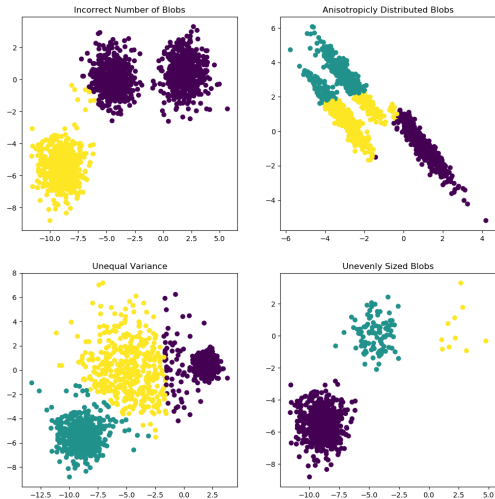
Then iterate these two steps until convergence of the centroids:

- 1 Assign each data point to its nearest centroid  
( $\rightarrow$  approximations necessary for high dimensions!)
- 2 Create  $k$  new centroids by taking the mean value of all of the data points assigned to each novel centroid  
( $\rightarrow$  approximations necessary for many data points)



# K-Means Clustering

Problematic data sets for k-means:



# K-Means Clustering

## Pros:

- conceptually simple algorithm
- mini-batches and different kind of initialization strategies are available (→ `k-means++`)
- scales to many data points (if approximations are utilized)

## Cons:

- sensitive to initialization of centroids
- can not model noise or outliers
- concave cluster shapes are problematic
- can not deal with uneven variance between clusters
- number  $k$  of clusters needs to be provided





# Lecture Overview

- 1 Motivation
- 2 Criteria for Clustering
- 3 K-Means
- 4 DBSCAN**

# DBSCAN

- DBSCAN (Ester et al., 1996) is a density-based, non-parametric clustering method.
- It received the **test of time award** at KDD conference in 2014.




# DBSCAN: A Cluster has High Density

# DBSCAN: A Cluster has High Density

- A **core point**  $p$  is a data point in an area of high density. It is defined by having at least **minPts** points (including  $p$ ) within an **eps**-neighborhood.
- A point  $q$  is **directly reachable** from  $p$ , if  $p$  is in the **eps** neighborhood of a core point  $q$ .
- A point  $q$  is **reachable** from  $p$  if there is a path along the points  $p_1, \dots, p_n$  with  $p_1 = p$  and  $p_n = q$ , where each  $p_{i+1}$  is directly reachable from  $p_i$ . Note that this implies that all points on the path must be core points, with the possible exception of  $q$  (if  $q$  is on the fringe of a cluster).
- All points not reachable from any other point are considered **outliers** or **noise** points.

If  $p$  is a core point, then it forms a cluster together with all points (core or non-core) that are reachable from it.

# DBSCAN



By Chire - Own work, CC BY-SA 3.0, <https://commons.wikimedia.org/w/index.php?curid=17045963>

Example with  $\text{minPts}=4$ . The  $\text{eps}$ -neighborhoods are indicated by circles.

# DBSCAN Algorithm

DBSCAN creates clusters by sequentially considering the training data points, starting with an arbitrary point:

- 1 Find the points in the **eps** neighborhood of every point, and identify the core points with more than **minPts** neighbors.
- 2 Find the connected components of core points on the neighbor graph, ignoring all non-core points.
- 3 Assign each non-core point to a nearby cluster if the cluster is an **eps** neighbor, otherwise assign it to noise.

# Pseudocode for DBSCAN

```
DBSCAN(DB, distFunc, eps, minPts) {  
    C = 0 /* Cluster counter */  
    for each point P in database DB {  
        if label(P)  $\neq$  undefined then continue /* Previously processed in inner loop */  
        Neighbors N = RangeQuery(DB, distFunc, P, eps) /* Find neighbors */  
        if |N| < minPts then { /* Density check */  
            label(P) = Noise /* Label as Noise */  
            continue  
        }  
        C = C + 1 /* next cluster label */  
        label(P) = C /* Label initial point */  
        Seed set S = N \ {P} /* Neighbors to expand */  
        for each point Q in S { /* Process every seed point */  
            if label(Q) = Noise then label(Q) = C /* Change Noise to border point */  
            if label(Q)  $\neq$  undefined then continue /* Previously processed */  
            label(Q) = C /* Label neighbor */  
            Neighbors N = RangeQuery(DB, distFunc, Q, eps) /* Find neighbors */  
            if |N|  $\geq$  minPts then { /* Density check */  
                S = S  $\cup$  N /* Add new neighbors to seed set */  
            }  
        }  
    }  
}  
  
RangeQuery(DB, distFunc, Q, eps) {  
    Neighbors = empty list  
    for each point P in database DB { /* Scan all points in the database */  
        if distFunc(Q, P)  $\leq$  eps then { /* Compute distance and check epsilon */  
            Neighbors = Neighbors  $\cup$  {P} /* Add to result */  
        }  
    }  
    return Neighbors  
}
```

# Characteristics of DBSCAN

## Pros:




- DBSCAN is fast! With  $n$  data points:  $O(n^2)$  as worst case (all points belong to single cluster), but  $O(n \log(n))$  with good data structures and for typical data.
- DBSCAN is deterministic for a fixed processing sequence.
- Number of clusters is determined automatically.
- Hyperparameter `min samples` can express prior knowledge about noise.
- Different distance metrics can be utilized.
- Hierarchical variant HDBSCAN is available.

## Cons:


- Varying the sequence of data points processed can lead to different clusterings.
- Hyperparameter `eps` is critical, no good default!



# Comparison of K-Means and DBSCAN

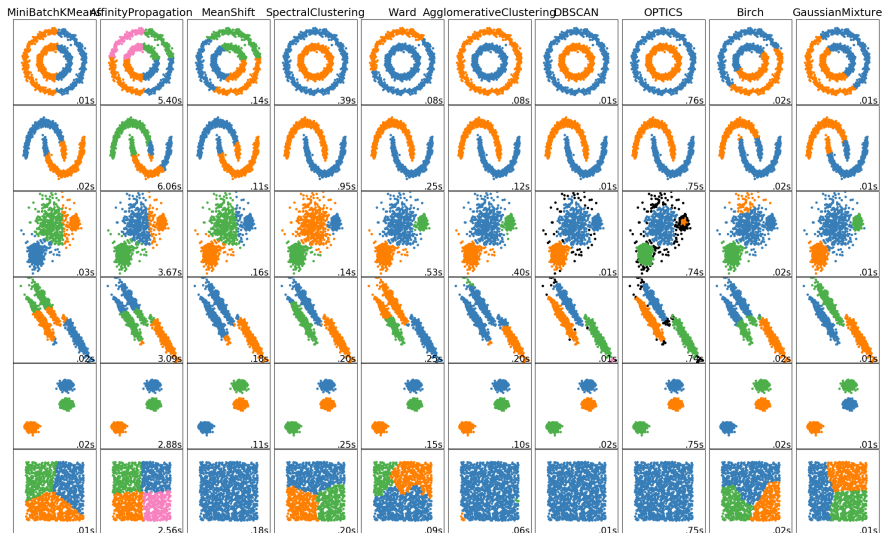


# A few Clustering Algorithms on Toy Data




<https://scikit-learn.org/stable/modules/clustering.html>

# A few Clustering Algorithms on Toy Data



<https://scikit-learn.org/stable/modules/clustering.html>

# Criteria for the Choice of Clustering Algorithms



Does the algorithm...

- expects each cluster to follow a specific distribution? (e.g. Gaussian)?
- considers density of data points?
- deal well with noisy data / high-dimensional data / redundant dimensions / irrelevant dimensions?
- deliver hard / soft clustering?
- deliver a strict partitioning (i.e. each object belongs to exactly one cluster)?
- deliver a hierarchical clustering?

# Wrap-Up: Summary by Learning Goals

Having heard this lecture and doing the assignment on clustering, you will be able to:

- Explain, which metrics can be used to create a clustering from unlabeled data
- formulate the optimization problem for k-means clustering and implement an iterative heuristic
- Describe pros and cons of k-means and DBSCAN
- Derive a metric for the quality of a given clustering (e.g. via the "**silhouette score**", see assignment)